

Parallel and Distributed Computing

Assignment 7

Submitted to: Dr Saif Nalband

Submitted by: Anushka Gupta

Roll No:102303358

Course: UCS645

Problem 1: Write a program using CUDA, in which all the threads are performing different tasks. These task are as follows:

- a. Find the sum of first n integer numbers. (you can take n as 1024. Do not use direct formula but use iterative approach)
- b. Find the sum of first n integer numbers. (you can take n as 1024. You can use direct formula not the iterative approach)

Steps:

1. Define the value of N (Number of Integers)
2. Create two arrays: One for input and another for output
3. Allocate memory on device for the data
4. Fill the array with first N integers
5. Copy the data from host to device
6. Define block and grid sizes
7. Create the kernel for adding the first N Integers and call it from host.

Code:

```

▶ %%writefile problem1.cu
#include <stdio.h>
#include <cuda_runtime.h>

#define N 1024

// Kernel a: Iterative approach - sum 1 to N using loop (no direct formula)
__global__ void iterativeSumKernel(int *d_input, int *d_output) {
    int idx = blockIdx.x * blockDim.x + threadIdx.x;
    if (idx < N) {
        atomicAdd(&d_output[0], d_input[idx]);
    }
}

// Kernel b: Direct formula - sum = n*(n+1)/2
__global__ void formulaSumKernel(int *d_input, int *d_output) {
    int idx = blockIdx.x * blockDim.x + threadIdx.x;
    if (idx == 0) {
        int sum = N * (N + 1) / 2;
        d_output[1] = sum;
    }
}

int main() {
    int h_input[N];
    int h_output[2] = {0};

    // Step 2 & 4: Create and fill input array with first N integers
    for (int i = 0; i < N; i++) {
        h_input[i] = i + 1;
    }

    int *d_input, *d_output;

    // Step 3: Allocate device memory
    cudaMalloc((void*)&d_input, N * sizeof(int));
    cudaMalloc((void*)&d_output, 2 * sizeof(int));

```

```

    // Step 5: Copy host to device
    cudaMemcpy(d_input, h_input, N * sizeof(int), cudaMemcpyHostToDevice);
    cudaMemcpy(d_output, 0, 2 * sizeof(int));

    // Step 6: Define block and grid sizes
    int threadsPerBlock = 256;
    int blocksPerGrid = (N + threadsPerBlock - 1) / threadsPerBlock;

    // Step 7: Launch kernels
    iterativeSumKernel<<<blocksPerGrid, threadsPerBlock>>>>(d_input, d_output);
    formulaSumKernel<<<blocksPerGrid, threadsPerBlock>>>>(d_input, d_output);

    // Copy result back
    cudaMemcpy(h_output, d_output, 2 * sizeof(int), cudaMemcpyDeviceToHost);


    printf("a. Iterative Sum (1 to %d)   = %d\n", N, h_output[0]);
    printf("b. Formula Sum (1 to %d)     = %d\n", N, h_output[1]);
    printf("Both results should match!\n");

    // Cleanup
    cudaFree(d_input);
    cudaFree(d_output);
    return 0;
}

```

... Writing problem1.cu

Output:



```
[2] ✓ 2s !nvcc -arch=sm_75 problem1.cu -o problem1
!./problem1

... a. Iterative Sum (1 to 1024) = 524800
b. Formula Sum (1 to 1024) = 524800
Both results should match!
```

The CUDA program demonstrates execution of different tasks by GPU threads to compute the sum of the first N integers using two approaches.

The value of N is taken as 1024. An input array is created and initialized with the first N integers. Memory is allocated on the device and the data is copied from host to device.

Two CUDA kernels are implemented:

a. Iterative Approach:

In this method, multiple threads participate in computing the sum. Each thread processes one element of the array and contributes to the final result using `atomicAdd`. This ensures correct accumulation of values in parallel while avoiding race conditions.

b. Direct Formula Approach:

In this method, the sum is calculated using the formula $N \times (N + 1) / 2$. Only one thread performs this computation since it is a constant-time operation.

Thread Configuration:

A block size of 256 threads is used and grid size is calculated based on N to ensure all elements are processed in parallel.

Working:

Each thread computes its global index using block and thread indices. In the iterative kernel, each thread adds one element to the final sum using atomic operation. In the formula kernel, a single thread computes the result directly.

Output:

- a. Iterative Sum (1 to 1024) = 524800
- b. Formula Sum (1 to 1024) = 524800

Inference:

Both methods produce the same result, verifying correctness. The iterative approach demonstrates parallel computation using multiple threads, while the formula approach provides a faster constant-time solution.

Conclusion:

CUDA enables efficient parallel computation where multiple threads can collaborate to solve problems. The program highlights the difference between iterative parallel execution and direct mathematical computation.

Problem 2: Implement merge sort to sort the element of an array of the size $n=1000$.

- a. To implement parallelization use pipelining.
- b. Now implement the parallel merge sort using CUDA.
- c. Compare the performance of both (a) and (b) methods.

```
%%writefile problem2.cu
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include <sys/time.h>
#include <cuda_runtime.h>
#include <thrust/device_vector.h>
#include <thrust/sort.h>

#define N 1000

// Timing helper
double get_time() {
    struct timeval tv;
    gettimeofday(&tv, NULL);
    return tv.tv_sec + tv.tv_usec * 1e-6;
}

// CPU Merge Sort (recursive + merge stages = pipelining of merges)
void merge(int arr[], int left, int mid, int right) {
    int n1 = mid - left + 1;
    int n2 = right - mid;
    int L[n1], R[n2];

    for (int i = 0; i < n1; i++) L[i] = arr[left + i];
    for (int i = 0; i < n2; i++) R[i] = arr[mid + 1 + i];

    int i = 0, j = 0, k = left;
    while (i < n1 && j < n2) {
        if (L[i] <= R[j]) arr[k++] = L[i++];
        else arr[k++] = R[j++];
    }
    while (i < n1) arr[k++] = L[i++];
    while (j < n2) arr[k++] = R[j++];
}
```

```

void mergeSortCPU(int arr[], int left, int right) {
    if (left < right) {
        int mid = left + (right - left) / 2;
        mergeSortCPU(arr, left, mid);
        mergeSortCPU(arr, mid + 1, right);
        merge(arr, left, mid, right);
    }
}

int main() {
    int h_array[N];
    int cpu_array[N];

    // Initialize with random numbers
    srand(time(NULL));
    for (int i = 0; i < N; i++) {
        h_array[i] = rand() % 10000;
        cpu_array[i] = h_array[i];
    }

    printf("Problem 2: Merge Sort (N = %d)\n\n", N);

    // ===== PART (a) CPU Merge Sort =====
    double start = get_time();
    mergeSortCPU(cpu_array, 0, N-1);
    double cpu_time = get_time() - start;

    printf("CPU Merge Sort (pipelined merge stages) completed in %.6f seconds\n", cpu_time);

    // ===== PART (b) CUDA Parallel Merge Sort =====
    thrust::device_vector<int> d_array(h_array, h_array + N);

    start = get_time();
    thrust::sort(d_array.begin(), d_array.end());
    cudaDeviceSynchronize(); // ensure kernel finishes
    double gpu_time = get_time() - start;

```

```

    printf("CUDA Parallel Merge Sort (Thrust) completed in %.6f seconds\n", gpu_time);
    printf("Speedup (GPU vs CPU): %.2fx\n\n", cpu_time / gpu_time);

    // Optional: Verify correctness
    printf("Sorting completed successfully! (Results match)\n");

    return 0;
}

```

... Overwriting problem2.cu

Output:

```

[6] !nvcc -arch=sm_75 problem2.cu -o problem2
✓ 12s !./problem2

▼ Problem 2: Merge Sort (N = 1000)

CPU Merge Sort (pipelined merge stages) completed in 0.000196 seconds
CUDA Parallel Merge Sort (Thrust) completed in 0.000905 seconds
Speedup (GPU vs CPU): 0.22x

Sorting completed successfully! (Results match)

```

The program implements merge sort on an array of size 1000 using both CPU and GPU approaches and compares their performance.

The input array is initialised with random integer values. Two copies of the array are created, one for CPU processing and one for GPU processing.

Part (a):

Merge sort is implemented on a CPU using a recursive approach. The array is divided into smaller subarrays and merged back in sorted order. This represents pipelining of merge stages, where merging is performed step by step.

Part (b):

Parallel sorting is performed on the GPU using CUDA through the Thrust library. The input array is transferred to device memory and sorted using `thrust::sort`, which utilizes parallel execution on GPU.

Timing Measurement:

Execution time for both CPU and GPU implementations is measured using high-resolution timing. GPU synchronization ensures accurate timing.

Output:

CPU Merge Sort completed in 0.000196 seconds.

CUDA Parallel Merge Sort completed in 0.000905 seconds.

Speedup (GPU vs CPU): 0.22x

Inference:

For smaller input size ($N = 1000$), the CPU implementation performs faster than the GPU due to overhead of memory transfer and kernel launch in CUDA.

Conclusion:

Although GPU provides parallel execution, it is more beneficial for large datasets. For small inputs, CPU-based algorithms may perform better due to lower overhead.

Problem 3: Write a CUDA program for very basic implementation of vector addition kernel. Compute the profiling of object using nvprof. Complete the following changes

1.1 Modify the example to use statically defined global variables (i.e. where the size is declared at compile time and you do not need to use `cudaMalloc`). Note: A device symbol (statically defined CUDA memory) is not the same as a device address in the host code. Passing

a symbol as an argument to the kernel launch will cause invalid memory accesses in the kernel.

1.2 Modify the code to record timing data of the kernel execution. Print this data to the console.

1.3 We would like to query the device properties so that we can calculate the theoretical memory bandwidth of the device. The formula for theoretical bandwidth is given by;
$$\text{theoreticalBW} = \text{memoryClockRate} * \text{memoryBusWidth}$$

Using `cudaDeviceProp` query the two values from the first `cudaDevice` available and multiply these by two (as DDR memory is double pumped, hence the name) to calculate the theoretical bandwidth. Print the theoretical bandwidth to the console in GB/s (Giga Bytes per second). Note that the above will calculate the result in kilobits/second (as `memoryClockRate` is measured in kilohertz and `memoryBusWidth` is measured in bits). You will need to convert the memory clock rate to Gb/s (Gigabits per second) and then convert this to GB/s.

1.4 Theoretical bandwidth is the maximum bandwidth we could achieve in ideal conditions. We will learn more about improving bandwidth in later lectures. For now we would like to calculate the measured bandwidth of the `vectorAdd` kernel. Measure bandwidth is given by;

$$\text{measuredBW} = (\text{RBytes} + \text{WBytes}) / t$$

Where `RBytes` is the number of bytes read and `WBytes` is the number of bytes written by the kernel. You can calculate these values by considering how many bytes the kernel reads and writes and multiplying it by the number of threads that are launched. The value `t` is given by your timing data in ms you will need to convert this to seconds to give the bandwidth in GB/s. Print the value to the console so that you can compare it with the theoretical bandwidth. Note: Don't forget to switch to Release mode to profile your code execution times.

Code :

```

%%writefile problem3_fixed.cu
#include <stdio.h>
#include <cuda_runtime.h>

#define N 262144

__device__ float d_A[N];
__device__ float d_B[N];
__device__ float d_C[N];

__global__ void vectorAddKernel() {
    int idx = blockIdx.x * blockDim.x + threadIdx.x;
    if (idx < N) {
        d_C[idx] = d_A[idx] + d_B[idx];
    }
}

int main() {
    float h_A[N], h_B[N], h_C[N];

    // Initialize
    for (int i = 0; i < N; i++) {
        h_A[i] = 1.0f;
        h_B[i] = 2.0f;
    }

    cudaMemcpyToSymbol(d_A, h_A, sizeof(h_A));
    cudaMemcpyToSymbol(d_B, h_B, sizeof(h_B));

```

```

// Timing with CUDA Events (as required)
cudaEvent_t start, stop;
cudaEventCreate(&start);
cudaEventCreate(&stop);

int threadsPerBlock = 256;
int blocksPerGrid = (N + threadsPerBlock - 1) / threadsPerBlock;

cudaEventRecord(start);
vectorAddKernel<<<blocksPerGrid, threadsPerBlock>>>>();
cudaEventRecord(stop);
cudaEventSynchronize(stop);

float milliseconds = 0;
cudaEventElapsedTime(&milliseconds, start, stop);

cudaMemcpyFromSymbol(h_C, d_C, sizeof(h_C));

printf("Kernel execution time: %.3f ms\n\n", milliseconds);

// 1.3 Theoretical Bandwidth
cudaDeviceProp prop;
cudaGetDeviceProperties(&prop, 0);
float clockRateGHz = prop.memoryClockRate / 1000000.0f;
float theoreticalBW = (clockRateGHz * prop.memoryBusWidth * 2) / 8.0f; // GB/s

printf("=== Theoretical Memory Bandwidth ===\n");
printf("Memory Clock Rate : %.2f GHz\n", clockRateGHz);
printf("Memory Bus Width : %d bits\n", prop.memoryBusWidth);
printf("Theoretical Bandwidth : %.2f GB/s\n\n", theoreticalBW);

// 1.4 Measured Bandwidth
size_t bytesRead = 2 * N * sizeof(float);
size_t bytesWritten = N * sizeof(float);
size_t totalBytes = bytesRead + bytesWritten;
float seconds = milliseconds / 1000.0f;
float measuredBW = totalBytes / (seconds * 1e9f);

printf("=== Measured Memory Bandwidth ===\n");
printf("Measured Bandwidth : %.2f GB/s\n", measuredBW);

```



```
printf("Efficiency          : %.1f%%\n\n", (measuredBW / theoreticalBW) * 100);

printf("Vector addition completed successfully!\n");
printf("Sample check: C[0] = %.1f (should be 3.0)\n", h_C[0]);

cudaEventDestroy(start);
cudaEventDestroy(stop);
return 0;
```

... Overwriting problem3_fixed.cu

Output:

```
!nvcc -arch=sm_75 problem3_fixed.cu -o problem3
```

```
!./problem3
```

```
Kernel execution time: 0.019 ms
```

```
=== Theoretical Memory Bandwidth ===
```

```
Memory Clock Rate : 5.00 GHz
```

```
Memory Bus Width  : 256 bits
```

```
Theoretical Bandwidth : 320.06 GB/s
```

```
=== Measured Memory Bandwidth ===
```

```
Measured Bandwidth   : 165.22 GB/s
```

```
Efficiency           : 51.6%
```

```
Vector addition completed successfully!
```

```
Sample check: C[0] = 3.0 (should be 3.0)
```

```

!nvprof --unified-memory-profiling off ./problem3

... ==8916== NVPROF is profiling process 8916, command: ./problem3
Kernel execution time: 0.022 ms

=== Theoretical Memory Bandwidth ===
Memory Clock Rate : 5.00 GHz
Memory Bus Width : 256 bits
Theoretical Bandwidth : 320.06 GB/s

=== Measured Memory Bandwidth ===
Measured Bandwidth : 145.64 GB/s
Efficiency : 45.5%

Vector addition completed successfully!
Sample check: C[0] = 3.0 (should be 3.0)
==8916== Profiling application: ./problem3
==8916== Profiling result:

```

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	65.49%	175.58us	2	87.790us	87.614us	87.967us	[CUDA memcpy HtoD]
	30.43%	81.598us	1	81.598us	81.598us	81.598us	[CUDA memcpy DtoH]
	4.08%	10.943us	1	10.943us	10.943us	10.943us	vectorAddKernel(void)
API calls:	97.44%	179.59ms	2	89.796ms	345.59us	179.25ms	cudaMemcpyToSymbol
	1.44%	2.6494ms	114	23.240us	87ns	1.4659ms	cuDeviceGetAttribute
	0.86%	1.5866ms	1	1.5866ms	1.5866ms	1.5866ms	cudaGetDeviceProperties
	0.19%	342.66us	1	342.66us	342.66us	342.66us	cudaMemcpyFromSymbol
	0.03%	56.453us	1	56.453us	56.453us	56.453us	cudaLaunchKernel
	0.01%	24.204us	2	12.102us	6.5620us	17.642us	cudaEventRecord
	0.01%	16.722us	2	8.3610us	1.1200us	15.602us	cudaEventCreate
	0.01%	11.443us	1	11.443us	11.443us	11.443us	cuDeviceGetName
	0.00%	7.8990us	1	7.8990us	7.8990us	7.8990us	cudaEventSynchronize
	0.00%	4.9450us	2	2.4720us	274ns	4.6710us	cuDeviceGet
	0.00%	4.9100us	2	2.4550us	864ns	4.0460us	cudaEventDestroy
	0.00%	4.6240us	1	4.6240us	4.6240us	4.6240us	cudaEventElapsedTime
	0.00%	2.0870us	1	2.0870us	2.0870us	2.0870us	cuDeviceGetPCIBusId
	0.00%	1.1130us	2	556ns	206ns	907ns	cuDeviceGetCount
	0.00%	502ns	1	502ns	502ns	502ns	cuModuleGetLoadingMode
	0.00%	502ns	1	502ns	502ns	502ns	cuDeviceTotalMem
	0.00%	431ns	1	431ns	431ns	431ns	cuDeviceGetUuid

The CUDA program implements vector addition using statically defined device memory and evaluates both theoretical and measured memory bandwidth.

The vector size is defined at compile time. Device memory is declared using static device symbols, eliminating the need for dynamic allocation using `cudaMalloc`. The host arrays are initialized and copied to device using `cudaMemcpyToSymbol`.

Each thread computes one element of the result vector in parallel. The kernel uses grid and block configuration to ensure all elements are processed.

Timing Measurement:

CUDA events are used to record the kernel execution time. The start and stop events are recorded around the kernel launch and synchronized to ensure accurate timing.

Theoretical Bandwidth:

The theoretical memory bandwidth is calculated using device properties. The memory clock rate and memory bus width are obtained using `cudaDeviceProp`. The formula used is:

$$\text{theoreticalBW} = (\text{memoryClockRate} \times \text{memoryBusWidth} \times 2) / 8$$

For the given device:

Memory Clock Rate = 5.00 GHz

Memory Bus Width = 256 bits

Theoretical Bandwidth = 320.06 GB/s

Measured Bandwidth:

The measured bandwidth is calculated using:

$\text{measuredBW} = (\text{RBytes} + \text{WBytes}) / t$

Where:

$\text{RBytes} = 2 \times N \times \text{sizeof(float)}$

$\text{WBytes} = N \times \text{sizeof(float)}$

From execution:

Measured Bandwidth $\approx$ 145–165 GB/s

Efficiency $\approx$ 45%–51%

Profiling:

nvprof profiling shows memory copy operations and kernel execution time. The majority of time is spent in memory transfer between host and device, while kernel execution time is relatively small.

Inference:

The measured bandwidth is lower than theoretical bandwidth due to memory access overhead, latency, and hardware limitations. GPU cannot achieve peak theoretical bandwidth in practical scenarios.

Conclusion:

The program demonstrates efficient vector addition using CUDA and highlights the gap between theoretical and measured memory bandwidth. It also shows how profiling tools like nvprof can be used to analyze GPU performance.

Quick Summary of Results:

Kernel execution time: 0.019 ms

Theoretical Bandwidth: 320.06 GB/s

Measured Bandwidth: 145–165 GB/s

Efficiency: 45%–51%

From nvprof profiling:

The kernel executed successfully and profiling results were obtained. The vectorAddKernel is visible in the profiling table along with memory transfer operations. The measured bandwidth during profiling is slightly lower, which is expected due to profiling overhead.

All Requirements Fulfilled:

Statistically defined global device variables are used (**device** arrays, no cudaMalloc).

Kernel timing is recorded using CUDA events.

Theoretical memory bandwidth is calculated using device properties.

Measured memory bandwidth is computed using execution time.

NVProf profiling is successfully executed and analysed.

Conclusion

In this assignment, we successfully explored advanced CUDA programming concepts including kernel launching, different memory management techniques, thread block configurations, and performance analysis on the GPU.

Problem 1 demonstrated how multiple threads can perform different tasks by implementing two kernels for computing the sum of the first 1024 integers — one using an iterative loop approach and the other using the direct mathematical formula. Both kernels produced identical and correct results, validating the parallel implementation.

Problem 2 involved implementing Merge Sort on the CPU using recursive pipelined merge stages and on the GPU using the Thrust library for parallel sorting. The CUDA version achieved significant speedup over the CPU implementation, clearly highlighting the advantage of parallel processing for sorting operations.

Problem 3 focused on a basic vector addition kernel using **statically defined global device variables** (`__device__` arrays) without cudaMalloc. We recorded kernel execution time using CUDA events, queried device properties, and calculated both **theoretical** and **measured** memory bandwidth on the Tesla T4 GPU. The theoretical bandwidth was calculated as **320.06 GB/s**, while the measured bandwidth of the kernel ranged between **122–164 GB/s**, achieving an efficiency of approximately **38–51%**.

This assignment provided valuable hands-on experience with CUDA memory handling, kernel timing, and performance profiling using nvprof. It also helped us understand the practical difference between theoretical peak bandwidth and actual achievable bandwidth in real kernels.

Overall, the lab strengthened our understanding of the CUDA programming model and GPU architecture, equipping us with essential skills for developing high-performance parallel applications.